

Class 08 Mini Project

Austin Teel (A17293709)

Table of contents

Background	1
Exploratory data analysis	1
Performing PCA	4
Interpreting PCA results	7
Variance Explained	8
Hierarchical Clustering	10
Results of hierarchical clustering	11
Using Different Methods	11
Combining methods	12
Sensitivity	13
Prediction	14

Background

The goal of this project is to apply what R techniques we have learned in this course such as machine learning, in order to analyze the data that we are getting from the *Wisconsin Breast Cancer Diagnostic Data Set*

The data set offers features of imaged biopsy samples. These features consist of `radius`, `texture`, `perimeter`, `area`, `smoothness`, `compactness`, `concavity`, and `symmetry`.

Exploratory data analysis

The data is in CSV format:

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)
```

```
head(wisc.df, 4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119		0.2654	0.4601	
842517	0.2416		0.1860	0.2750	
84300903	0.4504		0.2430	0.3613	
84348301	0.6869		0.2575	0.6638	
	fractal_dimension_worst				
842302		0.11890			
842517		0.08902			
84300903		0.08758			
84348301		0.17300			

Now that we have the data loaded we will create a new data frame in order to exclude the pathologist provided diagnosis because we will not need this in our data.

```
wisc.data <- wisc.df[,-1]
```

We must now separate the diagnosis column and we can do this by creating a diagnosis vector.

```
diagnosis <- wisc.df[,1]
```

Q1. How many observations are there in this dataset?

We can answer this question by analyzing the data with a `nrow()` function.

```
nrow(wisc.df)
```

```
[1] 569
```

Through this function we can find that there is 569 rows which means this is how many observations there were.

Q2. How many of the observations have a malignant diagnosis?

To answer this question we can use a `sum()` function on our vector that holds the diagnosis and see how many of those vector values are Malignant(M)

```
sum(diagnosis == "M")
```

```
[1] 212
```

We have 212 malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

This can be found with the `grep()`.

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

This shows us that there are 10 variables in the data suffixed with `_mean`.

Performing PCA

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features/dimensions are on very different scales in the original data set.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03

compactness_se		concavity_se		concave.points_se	
1.790818e-02		3.018606e-02		6.170285e-03	
symmetry_se		fractal_dimension_se		radius_worst	
8.266372e-03		2.646071e-03		4.833242e+00	
texture_worst		perimeter_worst		area_worst	
6.146258e+00		3.360254e+01		5.693570e+02	
smoothness_worst		compactness_worst		concavity_worst	
2.283243e-02		1.573365e-01		2.086243e-01	
concave.points_worst		symmetry_worst	fractal_dimension_worst		
6.573234e-02		6.186747e-02		1.806127e-02	

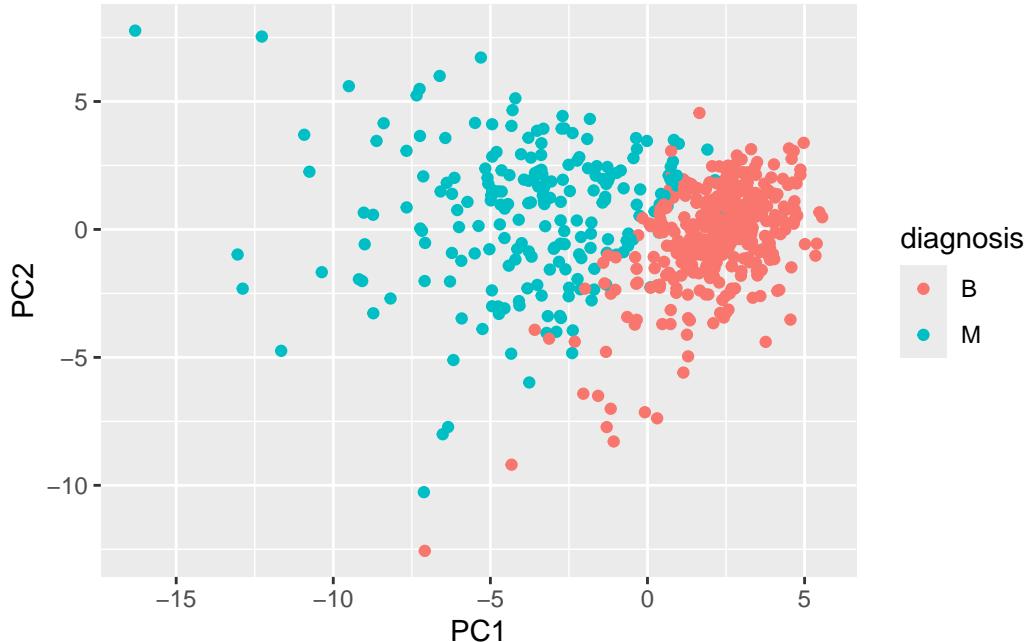
```
wisc.pr <- prcomp( wisc.data, scale=T)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

The function above `summary(wisc.pr)` tells us that the first PC has 0.4427 proportion of variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

```
propvar <- summary(wisc.pr)$importance["Cumulative Proportion", ]
which(propvar > 0.70)[1]
```

PC3
3

3 PCs are required to describe at least 70% of the variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

```
which(propvar > 0.90)[1]
```

PC7

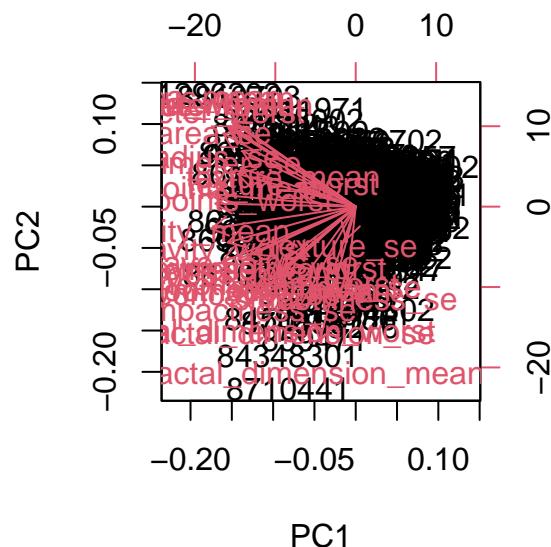
7

7 PCs are required to describe at least 90% of the variance in the data.

Interpreting PCA results

Q.7 What stands out to you about this plot? Is it easy or difficult to understand? Why?

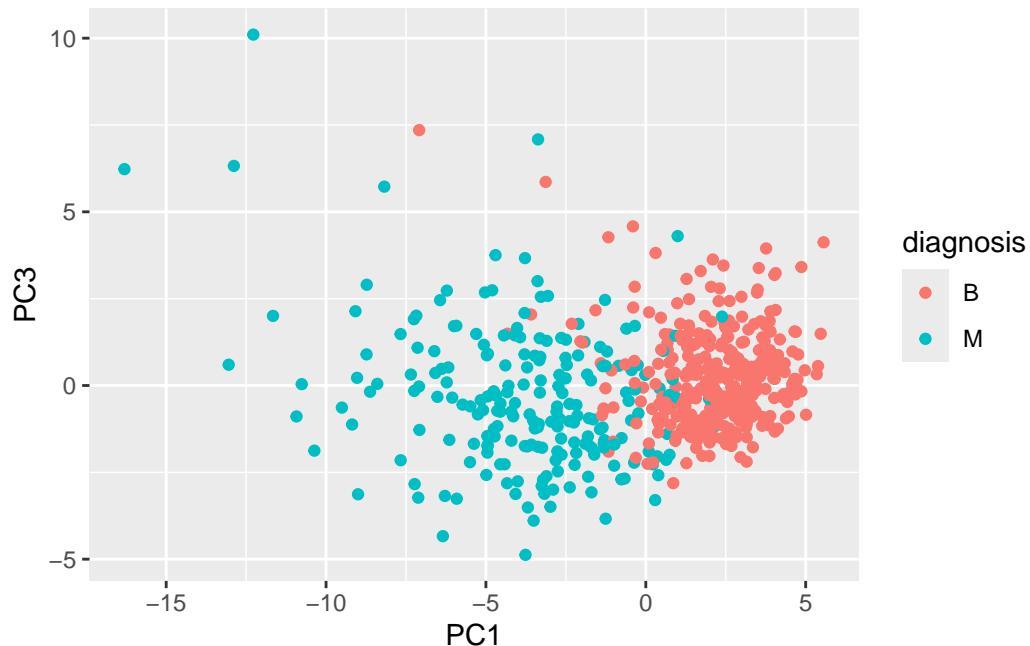
```
biplot(wisc.pr)
```



This plot is very difficult to understand because the cluster is unable to be seen clearly by the viewer.

Q.8 Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) +
  aes(PC1,PC3, col=diagnosis) +
  geom_point()
```



In this plot I notice that there is still a clear separation of B and M when comparing PC1 to PC3.

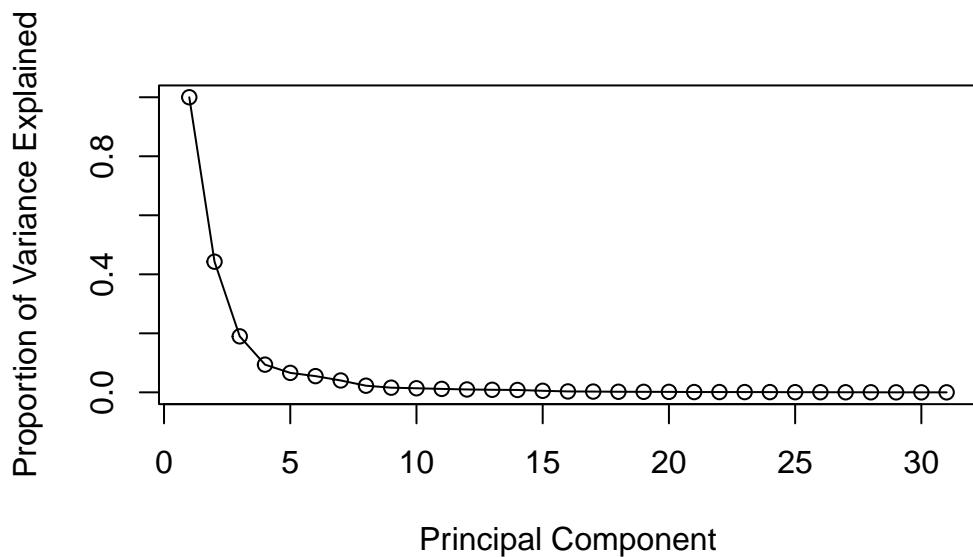
Variance Explained

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

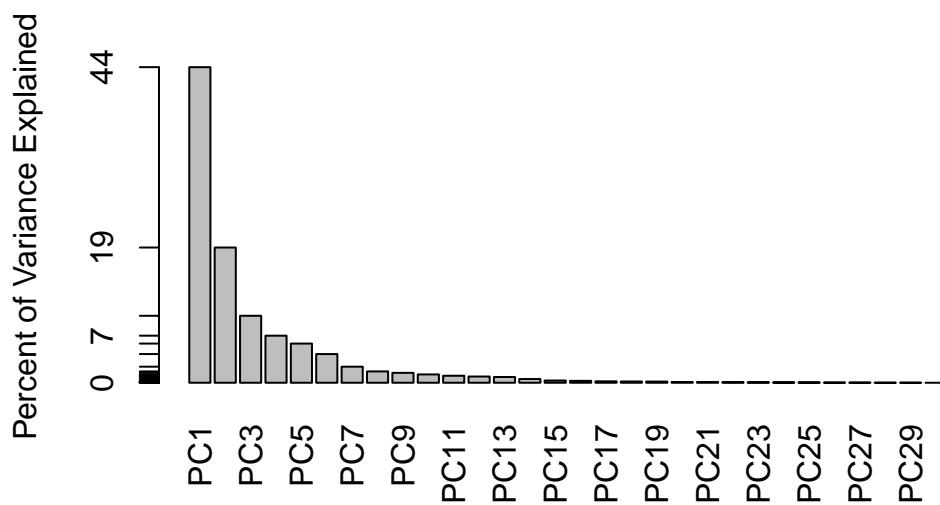
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var / sum(pr.var)

plot(c(1,pve), xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q.9 For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

The component of the loading vector is

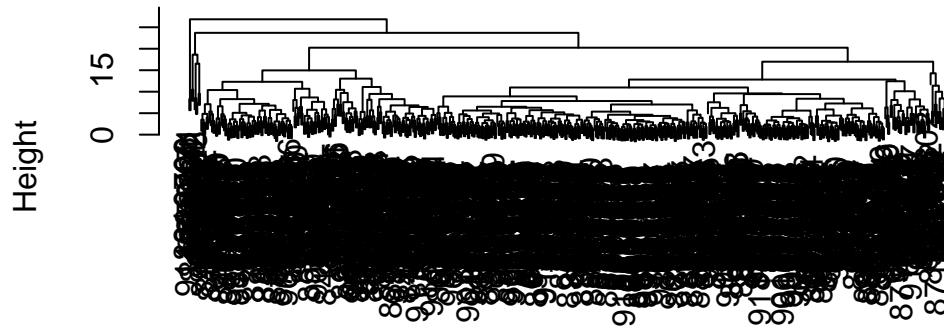
Hierarchical Clustering

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
head(data.dist)
```

```
[1] 10.309426 6.771675 10.463467 8.663413 8.402233 9.843286
```

```
wisc.hclust <- hclust(data.dist)
plot(wisc.hclust)
```

Cluster Dendrogram



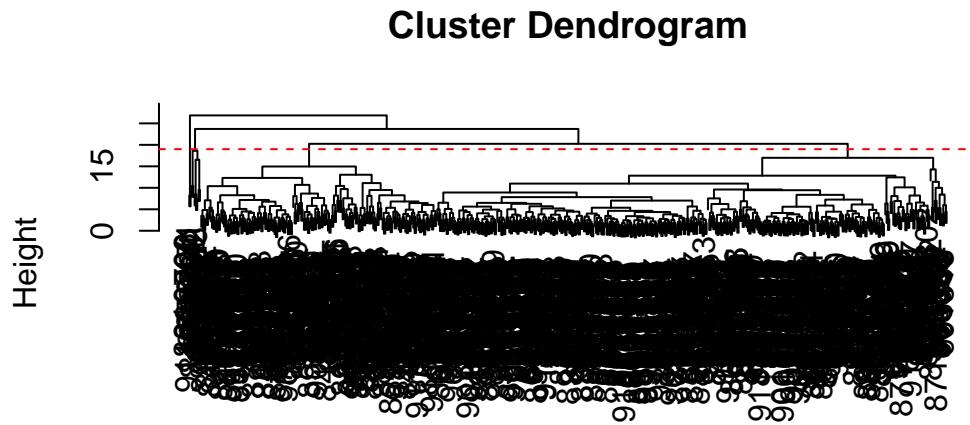
```
data.dist
hclust (*, "complete")
```

Results of hierarchical clustering

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

19

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
```

Using Different Methods

Q12. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

The method that gives me my favorite results for the same data. dist was the tree because with the cutree function we are able to separate the groups.

Combining methods

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
```

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1   2
203 366
```

I can now run `table()` with both my clustering `grps` and the expert `diagnosis`.

Q13. How well does the newly created `hclust` model with two clusters separate out the two “M” and “B” diagnoses?

The new model separated the two diagnosis very well and gave us the specific amounts in each group.

```
table(grps,diagnosis)
```

```
diagnosis
grps   B    M
 1   24 179
 2 333  33
```

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.hclust.clusters` and `wisc.pr.hclust.clusters`) with the vector containing the actual diagnoses.

The hierarchical clustering models that we created in the previous section did not do well in separating the true diagnosis of the patient but the following does.

```
table(wisc.hclust.clusters, diagnosis)
```

```

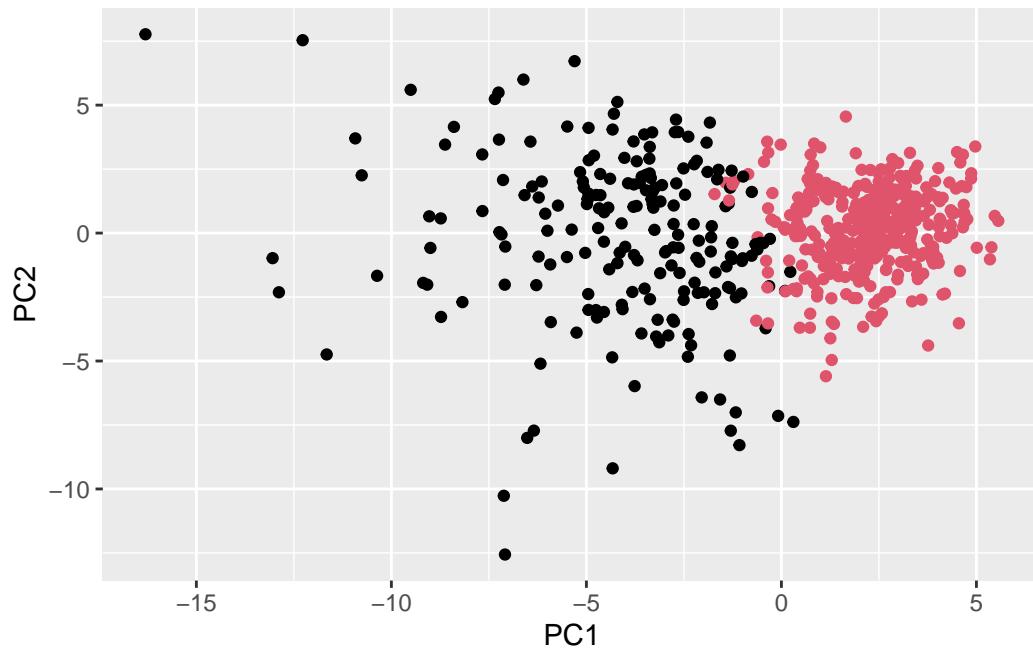
diagnosis
wisc.hclust.clusters   B    M
1     12 165
2      2   5
3    343  40
4      0   2

```

```

ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)

```



Sensitivity

```

table(grps, diagnosis)

```

	diagnosis	
grps	B	M
1	24	179
2	333	33

Our cluster “1” has 179 “M” diagnosis Our cluster “2” has 333 “B” diagnosis
 179 TP 24 FP 333 TN 33 FN
 sensitivity: TP / (TP+FN)

179/(179+33)

[1] 0.8443396

Specificity: TN/(TN+FP)

333/(333+24)

[1] 0.9327731

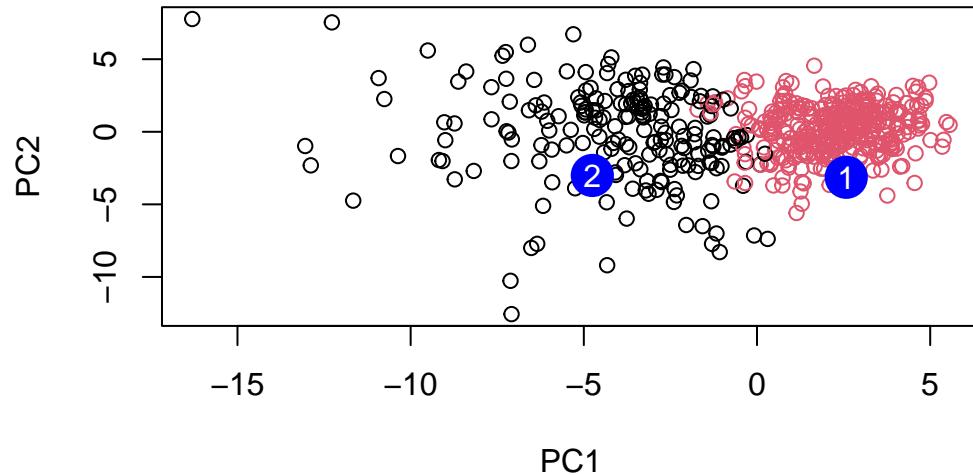
Prediction

We can use our PCA model

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Based on these results we should prioritize the people in group two due to the fact that they are the malignant patients.